

Programming Languages

Multiple Inheritance

Dr. Michael Petter
Winter Term 2021

Outline

Inheritance Principles

- 1 Interface Inheritance
- 2 Implementation Inheritance
- 3 Dispatching implementation choices

C++ Object Heap Layout

- 1 Basics
- 2 Single-Inheritance
- 3 Virtual Methods

C++ Multiple Parents Heap Layout


- 1 Multiple-Inheritance
- 2 Virtual Methods
- 3 Common Parents

Excursion: Linearization

- 1 Ambiguous common parents
- 2 Principles of Linearization
- 3 Linearization algorithms

“Wouldn’t it be nice to inherit from several parents?”

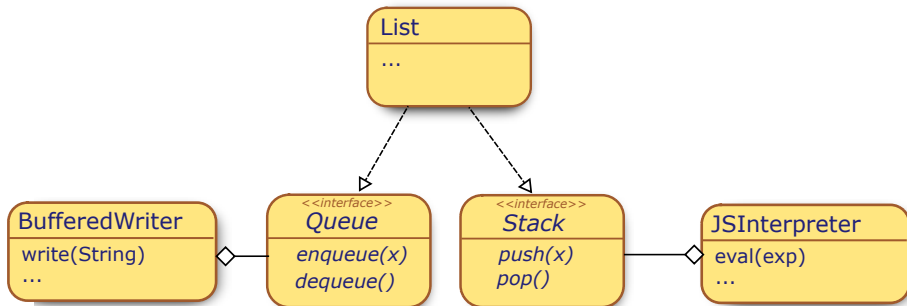
Implementation inheritance



The classic motivation for inheritance is implementation inheritance

- *Code reuse*
- Parent acts as library of common behaviours
- Child specializes parents, replacing particular methods with custom ones

Interface Inheritance



Code sharing in interface inheritance inverts this relation

- *Behaviour contract*
- Parent formalizes type requirements to participate in shared code frames
- Child provides functionality in methods with signatures predefined by the parent

“So how do we lay out objects in memory anyway?”

Excursion: Brief introduction to LLVM IR

LLVM intermediate representation as reference semantics:

```
;(recursive) struct definitions
%struct.A = type { i32, %struct.B, i32(i32)* }
%struct.B = type { i64, [10 x [20 x i32]], i8 }

;(stack-) allocation of objects
%a = alloca %struct.A
;adress computation for selection in structure (pointers):
%1 = getelementptr %struct.A* %a, i64 0, i64 2
;load from memory
%2 = load i32(i32)* %1
;indirect call
%retval = call i32 (i32)* %2(i32 42)
```

Retrieve the memory layout of a compilation unit with:

```
clang -cc1 -x c++ -v -fdump-record-layouts -emit-llvm source.cpp
```

Retrieve the IR Code of a compilation unit with:


```
clang -O1 -S -emit-llvm source.cpp -fno-discard-value-names -o /dev/stdout | llvm-cxxfilt
```

Object layout

```
class A {                public:
    int a; int f(int);
};
class B : public A { public:
    int b; int g(int);
};
class C : public B { public:
    int c; int h(int);
};

...

C c;
c.g(42);
```




```
%class.C = type { %class.B, i32 }
%class.B = type { %class.A, i32 }
%class.A = type { i32 }
```

```
%c = alloca %class.C
%1 = bitcast %class.C* %c to %class.B*
%2 = call i32 @_g(%class.B* %1, i32 42) ; g is statically known
```


Translation of a method body

```
class A {                public:
    int a; int f(int);
};
class B : public A { public:
    int b; int g(int);
};
class C : public B { public:
    int c; int h(int);
};
int B::g(int p) {
    return p+b;
};
```



```
%class.C = type { %class.B, i32 }
%class.B = type { %class.A, i32 }
%class.A = type { i32 }
```

```
define i32 @_g(%class.B* %this, i32 %p) {
    %1 = getelementptr %class.B* %this, i64 0, i32 1
    %2 = load i32* %1
    %3 = add i32 %2, %p
    ret i32 %3
}
```

“Now what about polymorphic calls?”

Single-Dispatching implementation choices

Single-Dispatching needs runtime action:

- 1 Manual search run through the super-chain (Java Interpreter \rightsquigarrow last talk)

```
call i32 @__dispatch(%class.C* %c,i32 42,i32* "f(int,void)")
```

- 2 Caching the dispatch result (\rightsquigarrow Hotspot/JIT)

```
; caching the recent result value of the __dispatch function  
; call i32 @__dispatch(%class.C* %c,i32 42)  
assert (%c type %class.D) ; verify objects class presumption  
call i32 @_f_from_D(%class.C* %c, i32 42) ; directly call f
```

- 3 Precomputing the dispatching result in tables


- 1 Full 2-dim matrix
- 2 1-dim Row Displacement Dispatch Tables
- 3 Virtual Tables
(\rightsquigarrow LLVM/GNU C++,this talk)

A	B	F					
1	1	2	...	8	9		7

	f()	g()	h()	i()	j()	k()	l()	m()	n()
A	1								
B	1	2							
C	3		4						
D	3	2	4	5					
E						6		7	
F					8	9		7	

Object layout – virtual methods

```
class A { public:  
    int a; virtual int f(int);  
           virtual int g(int);  
           virtual int h(int);  
};  
class B : public A { public:  
    int b; int g(int);  
};  
class C : public B { public:  
    int c; int h(int);  
}; ...  
C c;  
c.g(42);
```




```
%class.C = type { %class.B, i32, [4 x i8] }  
%class.B = type { [12 x i8], i32 }  
%class.A = type { i32 (...)**, i32 }
```

```
%c.vptr = bitcast %class.C* %c to i32 (%class.B*, i32)*** ; vtbl  
%1 = load (%class.B*, i32)*** %c.vptr ; dereference vptr  
%2 = getelementptr %1, i64 1 ; select g()-entry  
%3 = load (%class.B*, i32)** %2 ; dereference g()-entry  
%4 = call i32 @%3(%class.B* %c, i32 42)
```


“So how do we include several parent objects?”

Multiple inheritance class diagram



Static Type Casts

```
class A { public:  
    int a; int f(int);  
};  
class B { public:  
    int b; int g(int);  
};  
class C : public A , public B {  
    public:  
    int c; int h(int);  
};  
...  
B* b = new C();
```



```
%class.C = type { %class.A, %class.B, i32 }  
%class.A = type { i32 }  
%class.B = type { i32 }
```


```
%1 = call i8* @_new(i64 12)  
call void @_memset.p0i8.i64(i8* %1, i8 0, i64 12, i32 4, i1 false)  
%2 = getelementptr i8* %1, i64 4 ; select B-offset in C  
%b = bitcast i8* %2 to %class.B*
```

⚠ implicit casts potentially add a constant to the object pointer.

⚠ `getelementptr` implements ΔB as $4 \cdot i8$!

Keeping Calling Conventions

```
class A { public:  
    int a; int f(int);  
};  
class B { public:  
    int b; int g(int);  
};  
class C : public A , public B {  
    public:  
    int c; int h(int);  
};  
...  
C c;  
c.g(42);
```



```
%class.C = type { %class.A, %class.B, i32 }  
%class.A = type { i32 }  
%class.B = type { i32 }
```

```
%c = alloca %class.C  
%1 = bitcast %class.C* %c to i8*  
%2 = getelementptr i8* %1, i64 4 ; select B-offset in C  
%3 = call i32 @_g(%class.B* %2, i32 42) ; g is statically known
```


Ambiguities

```
class A { public: void f(int); };  
class B { public: void f(int); };  
class C : public A, public B {};  
  
C* pc;  
pc->f(42);
```

⚠ Which method is called?

Solution I: Explicit qualification

```
pc->A::f(42);  
pc->B::f(42);
```

Solution II: Automagical resolution

Idea: The Compiler introduces a linear order on the nodes of the inheritance graph

Linearization

Principle 1: Inheritance Relation

Defined by parent-child. Example:

$$C(A, B) \implies C \multimap A \wedge C \multimap B$$



Principle 2: Multiplicity Relation

Defined by the succession of multiple parents. Example: $C(A, B) \implies A \rightarrow B$



In General:

- 1 Inheritance is a uniform mechanism, and its searches (\rightarrow total order) apply identically for all object fields or methods
- 2 In the literature, we also find the set of constraints to create a linearization as Method Resolution Order
- 3 Linearization is a best-effort approach at best

MRO via DFS

Leftmost Preorder Depth-First Search

$$L[A] = ABWC$$

⚠ Principle 1 *inheritance* is violated

Python: classical python objects (≤ 2.1) use LPDFS!

LPDFS with Duplicate Cancellation

$$L[A] = ABCW$$

✓ Principle 1 *inheritance* is fixed



Python: new python objects (2.2) use LPDFS(DC)!

LPDFS with Duplicate Cancellation

$$L[A] = ABCWV$$

⚠ Principle 2 *multiplicity* not fulfillable

⚠ However $B \rightarrow C \implies W \rightarrow V??$




MRO via Refined Postorder DFS

Reverse Postorder Rightmost DFS

$$L[A] = ABFDCEGHW$$

✓ Linear extension of inheritance relation



$A(B, C) \ B(F, D) \ C(E, H)$
 $D(G) \ E(G) \ F(W) \ G(W) \ H(W)$

RPRDFS

$$L[A] = ABCDGEF$$


⚠ But principle 2 *multiplicity* is violated!

CLOS: uses Refined RPDFS [3]

Refined RPRDFS

$$L[A] = ABCDEFG$$

✓ Refine graph with conflict edge & rerun RPRDFS!



$A(B, C) \ B(F, G) \ C(D, E)$
 $D(G) \ E(F)$

MRO via Refined Postorder DFS

Refined RPRDFS


⚠ *Monotonicity* is not guaranteed!

Extension Principle: Monotonicity

If $C_1 \rightarrow C_2$ in C 's linearization, then $C_1 \rightarrow C_2$ for every linearization of C 's children.

$$L[A] = A \ B \ C \ D \ E \ F \ G \implies F \rightarrow G$$

$$L[C] = C \ D \ G \ E \ F \implies G \rightarrow F$$



MRO via C3 Linearization

A linearization L is an attribute $L[C]$ of a class C . Classes B_1, \dots, B_n are superclasses to child class C , defined in the *local precedence order* $C(B_1 \dots B_n)$. Then


$$L[C] = C \cdot \bigsqcup (L[B_1], \dots, L[B_n], B_1 \cdot \dots \cdot B_n) \quad | \quad C(B_1, \dots, B_n)$$

with

$$\bigsqcup_i (L_i) = \begin{cases} c \cdot (\bigsqcup_i (L_i \setminus c)) & \text{if } \exists_{\min k} \forall_j c = \text{head}(L_k) \notin \text{tail}(L_j) \\ \triangle \text{ fail} & \text{else} \end{cases}$$

MRO via C3 Linearization

$L[G]$	G
$L[F]$	F
$L[E]$	$E \cdot F$
$L[D]$	$D \cdot G$
$L[B]$	$B \cdot F \cdot G$
$L[C]$	$C \cdot D \cdot G \cdot E \cdot F$
$L[A]$	⚠ fail



C3 detects and reports a violation of *monotonicity* with the addition of $A(B, C)$ to the class set.
C3 linearization [1]: is used in *Python 3*, *Raku*, and *Solidity*

Linearization vs. explicit qualification

Linearization

- No switch/duplexer code necessary
- No explicit naming of qualifiers
- Unique *super* reference
- Reduces number of multi-dispatching conflicts

Qualification

- More flexible, fine-grained
- Linearization choices may be awkward or unexpected


Languages with automatic linearization exist

- *CLOS* Common Lisp Object System
- *Solidity*, *Python 3* and *Raku* with C3
- Prerequisite for → Mixins

“And what about dynamic dispatching in Multiple Inheritance?”

Virtual Tables for Multiple Inheritance

```
class A {                                     public:
    int a; virtual int f(int);
};
class B {                                     public:
    int b; virtual int f(int);
           virtual int g(int);
};
class C : public A , public B { public:
    int c; virtual int f(int);
};
...
C c;
B* pb = &c;
pb->f(42);
```




```
%class.C = type { %class.A, [12 x i8], i32 }
%class.A = type { i32 (...)**, i32 }
%class.B = type { i32 (...)**, i32 }
```

```
; B* pb = &c;
%0 = bitcast %class.C* %c to i8*      ; type fumbling
%1 = getelementptr i8* %0, i64 16     ; offset of B in C
%2 = bitcast i8* %1 to %class.B*      ; get typing right
store %class.B* %2, %class.B** %pb    ; store to pb
```

Virtual Tables for Multiple Inheritance

```
class A {                                     public:
    int a; virtual int f(int);
};
class B {                                     public:
    int b; virtual int f(int);
           virtual int g(int);
};
class C : public A , public B { public:
    int c; virtual int f(int);
};
...
C c;
B* pb = &c;
pb->f(42);
```



```
%class.C = type { %class.A, [12 x i8], i32 }
%class.A = type { i32 (...)**, i32 }
%class.B = type { i32 (...)**, i32 }
```

```
; pb->f(42);
%0 = load %class.B** %pb
%1 = bitcast %class.B* %0 to i32 (%class.B*, i32)***
%2 = load i32(%class.B*, i32)*** %1
%3 = getelementptr i32 (%class.B*, i32)** %2, i64 0
%4 = load i32(%class.B*, i32)** %3
%5 = call i32 @f(%class.B* %0, i32 42)


;load the b-pointer
;cast to vtable
;load vp_ptr
;select f() entry
;load function pointer
```

Basic Virtual Tables (→ C++-ABI)

A Basic Virtual Table

consists of different parts:

- 1 *offset to top* of an enclosing objects memory representation
- 2 *typeinfo pointer* to an RTTI object (not relevant for us)
- 3 *virtual function pointers* for resolving virtual methods




- Virtual tables are composed when multiple inheritance is used
- The `vptr` fields in objects are pointers to their corresponding virtual-subtables
- Casting preserves the link between an object and its corresponding virtual-subtable
- `clang -cc1 -fdump-vtable-layouts -emit-llvm source.cpp` yields the vtables of a compilation unit

Casting Issues

```
class A { public: virtual int f(int); int a; };  
class B { public: virtual int f(int); };  
class C : public A , public B {  
    public: virtual int f(int); int c;  
};  
C* c = new C();  
c->f(42);
```

```
B* b = new C();  
b->f(42);
```



⚠ this-Pointer for C::f is expected to point to c

Thunks

Solution: **thunks**

...are trampoline methods, delegating the virtual method to its original implementation with an adapted this-reference

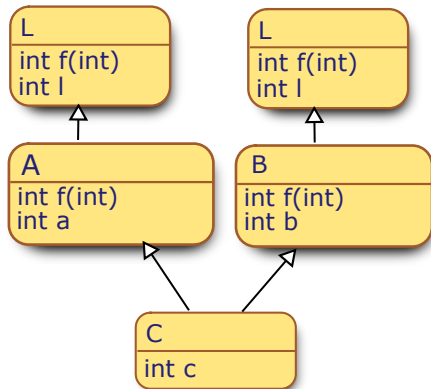
```
define i32 @__f(%class.B* %this, i32 %i) {  
    %1 = bitcast %class.B* %this to i8*  
    %2 = getelementptr i8* %1, i64 -16          ; sizeof(A)=16  
    %3 = bitcast i8* %2 to %class.C*  
    %4 = call i32 @__f(%class.C* %3, i32 %i)  
    ret i32 %4  
}
```

- ↪ B-in-C-vtable entry for `f(int)` is the thunk `_f(int)`
- ↪ `_f(int)` adds a compiletime constant ΔB to `this` before calling `f(int)`
- ↪ `f(int)` addresses its locals relative to what it assumes to be a C pointer

“But what if there are common ancestors?”


Common Bases – Duplicated Bases

Standard C++ multiple inheritance conceptually duplicates representations for common ancestors:



Duplicated Base Classes

```
class L {                                public:
    int l; virtual void f(int);
};
class A : public L {                    public:
    int a; void f(int);
};
class B : public L {                    public:
    int b; void f(int);
};
class C : public A , public B { public:
    int c;
};
...
C c;
L* p1 = (B*)&c;
p1->f(42); // where to dispatch?
C* pc = (C*)(B*)p1;
```




```
%class.C = type { %class.A, %class.B,
                  i32, [4 x i8] }
%class.A = type { [12 x i8], i32 }
%class.B = type { [12 x i8], i32 }
%class.L = type { i32 (...)**, i32 }
```

⚠ Ambiguity!

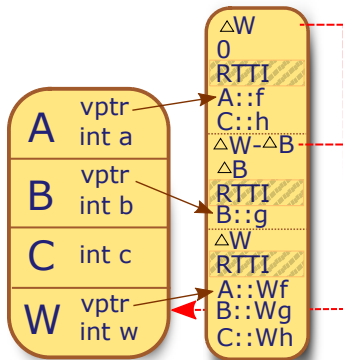
Common Bases – Shared Base Class

Optionally, C++ multiple inheritance enables a shared representation for common ancestors, creating the *diamond pattern*:



Shared Base Class

```
class W {                                     public:
    int w; virtual void f(int);
    virtual void g(int);
    virtual void h(int);
};
class A : public virtual W {                 public:
    int a; void f(int);
};
class B : public virtual W {                 public:
    int b; void g(int);
};
class C : public A, public B {               public:
    int c; void h(int);
};
...
C* pc;
pc->B::f(42);
((W*)pc)->h(42);
((B*)pc)->f(42);
```





⚠ Ambiguities

↪ e.g. overriding `f` in `A` and `B`

⚠ Offsets to virtual base

Dynamic Type Casts

```
class A : public virtual W {  
    ...  
};  
class B : public virtual W {  
    ...  
};  
class C : public A , public B {  
    ...  
};  
class D : public C,  
           public B {  
    ...  
};  
...  
C c;  
W* pw = &c;  
C* pc = dynamic_cast<C*>(pw);
```



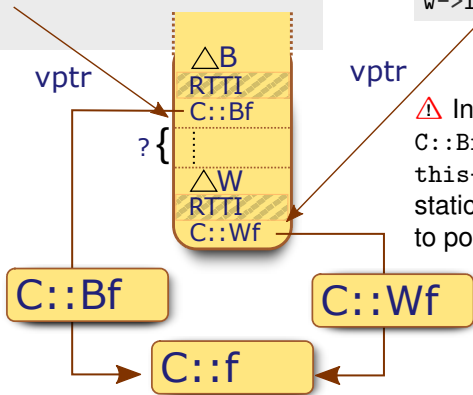
⚠ No guaranteed *constant* offsets between virtual bases and subclasses ~> No static casting!

⚠ *Dynamic casting* makes use of *offset-to-top*

Again: Casting Issues

```
class W { public: virtual int f(int); };  
class A : virtual W { int a; };  
class B : virtual W { int b; };  
class C : public A , public B {  
    public: virtual int f(int); int c;  
};  
B* b = new C();  
b->f(42);
```

```
W* w = new C();  
w->f(42);
```

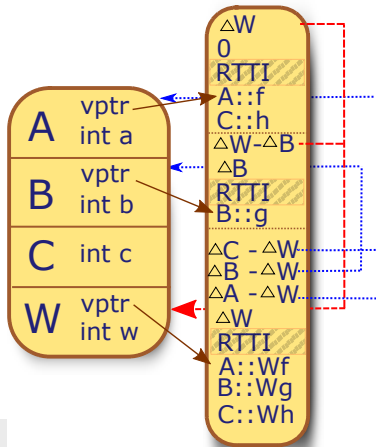


⚠ In a conventional thunk
C::Bf adjusts the
this-pointer with a
statically known constant
to point to C

Virtual Thunks

```
class W { ...  
    public: virtual void g(int);  
};  
class A : public virtual W {...};  
class B : public virtual W {  
    public: virtual void g(int i){ }; int b;  
};  
class C : public A,public B{...};  
C c;  
W* pw = &c;  
pw->g(42);
```

```
define void @_g(%class.W* %this, i32 %i) { ; virtual thunk to B::g  
    %1 = bitcast %class.W* %this to i8*  
    %2 = bitcast i8* %1 to i8**  
    %3 = load i8** %2 ; load W-vtable ptr  
    %4 = getelementptr i8* %3, i64 -32 ; -32 bytes is g-entry in vcall offsets  
    %5 = bitcast i8* %4 to i64*  
    %6 = load i64* %5 ; load g's vcall offset  
    %7 = getelementptr i8* %1, i64 %6 ; navigate to Wtop+vcalloffset  
    %8 = bitcast i8* %7 to %class.B*  
    call void @_g(%class.B* %8, i32 %i)  
    ret void  
}
```



Virtual Tables for Virtual Bases (\rightsquigarrow C++-ABI)


A Virtual Table for a Virtual Subclass

gets a *virtual base pointer*

A Virtual Table for a Virtual Base

consists of different parts:

- 1 *virtual call offsets* per virtual function for adjusting `this` dynamically
- 2 *offset to top* of an enclosing objects heap representation
- 3 *typeinfo pointer* to an RTTI object (not relevant for us)
- 4 *virtual function pointers* for resolving virtual methods



Virtual Base classes have *virtual thunks* which look up the offset to adjust the `this` pointer to the correct value in the virtual table!

Compiler and Runtime Collaboration

Compiler generates:

- 1 ... one code block for each method
- 2 ... one virtual table for each class-composition, with
 - ▶ references to the most recent implementations of methods of a *unique common signature* (\rightsquigarrow single dispatching)
 - ▶ sub-tables for the composed subclasses
 - ▶ static top-of-object and virtual bases offsets per sub-table
 - ▶ (virtual) thunks as `this`-adapters per method and subclass if needed

Runtime:

- 1 At program startup virtual tables are globally created
- 2 Allocation of memory space for each object followed by constructor calls
- 3 Constructor stores pointers to virtual table (or fragments) in the objects
- 4 Method calls transparently call methods statically or from virtual tables, *unaware of real class identity*
- 5 Dynamic casts may use *offset-to-top* field in objects

Polemics of Multiple Inheritance

Full Multiple Inheritance (FMI)

- Removes constraints on parents in inheritance
- More convenient and simple in the common cases
- Occurance of diamond pattern not as frequent as discussions indicate

Multiple Interface Inheritance (MII)

- simpler implementation
- Interfaces and aggregation already quite expressive
- Too frequent use of FMI considered as flaw in the class hierarchy design

Lessons Learned

Lessons Learned

- 1 Different purposes of inheritance
- 2 Heap Layouts of hierarchically constructed objects in C++
- 3 Virtual Table layout
- 4 LLVM IR representation of object access code
- 5 Linearization as alternative to explicit disambiguation
- 6 Pitfalls of Multiple Inheritance

Sidenote for MS VC++

- the presented approach is implemented in GNU C++ and LLVM
- Microsoft's MS VC++ approaches multiple inheritance differently
 - ▶ splits the virtual table into several smaller tables
 - ▶ keeps a vbptr (virtual base pointer) in the object representation, pointing to the virtual base of a subclass.

Further reading...

- [1] K. Barrett, B. Cassels, P. Haahr, D. Moon, K. Playford, and T. Withington.
A monotonic superclass linearization for dylan.
In *Object Oriented Programming Systems, Languages, and Applications*, 1996.
- [2] CodeSourcery, Compaq, EDG, HP, IBM, Intel, R. Hat, and SGI.
Itanium C++ ABI.
URL: <http://www.codesourcery.com/public/cxx-abi>.
- [3] R. Ducournau and M. Habib.
On some algorithms for multiple inheritance in object-oriented programming.
In *Proceedings of the European Conference on Object-Oriented Programming (ECOOP)*, 1987.
- [4] R. Kleckner.
Bringing clang and llvm to visual c++ users.
URL: <http://llvm.org/devmtg/2013-11/#talk11>.
- [5] B. Liskov.
Keynote address – data abstraction and hierarchy.
In *Addendum to the proceedings on Object-oriented programming systems, languages and applications*, OOPSLA '87, pages 17–34, 1987.
- [6] L. L. R. Manual.
Llvm project.
URL: <http://llvm.org/docs/LangRef.html>.
- [7] R. C. Martin.
The liskov substitution principle.
In *C++ Report*, 1996.
- [8] P. Sabanal and M. Yason.
Reversing c++.
In *Black Hat DC*, 2007.
URL: https://www.blackhat.com/presentations/bh-dc-07/Sabanal_Yason/Paper/bh-dc-07-Sabanal_Yason-WP.pdf.
- [9] B. Stroustrup.
Multiple inheritance for C++.
In *Computing Systems*, 1999.

Mini Seminars

- 1 SC=CC in Multicore Architectures with Cache (Meixner/Sorin 2006/2009)
- 2 Litmus Testing Memory Models: Herdtools 7
- 3 The Linux Kernel Memory Model
- 4 A Formal Analysis of the NVIDIA PTX Memory Consistency Model (2019)
- 5 GPU Concurrency: Weak Behaviours and Programming Assumptions (2015)
- 6 Transactional Memory Systems other than TSX: IBM Power 8 / BlueGene / zEnterprise
- 7 Lambda Calculus: Y Combinator and Recursion / SKI Combinator
- 8 Templates vs. Inheritance